

Reinforcement
Learning
Course
by
David Silver

Yury Gassimov
Notes

Sutton's

— 2023 —

21.01.2023,

L.1: Intro to Reinforcement Learning.

Books: Sutton & Barto. An introduction to reinforcement learning. 1998.


↗ Szepesvári. Algorithms for RL. 2010.
↖ more mathematical.

Example: • manage an investment portfolio

Reward R_t — scalar feedback signal.

Def. (Reward hypothesis) All goals can be described by the maximisation of expected cumulative reward.

Goal: select actions to max. total future reward
• You cannot use greedy algos.



The history is the sequence:

$$H_t = A_1, O_1, R_1, \dots, A_t, O_t, R_t$$

State is the information used to determine what happens next.

$$S_t = f(H_t)$$

Environment state S_t^e is the environment's private representation.
not usually visible to the agent.

Agent state S_t^a is the agent's internal representation.
 $S_t^a = f(H_t)$

Def. A state S_t is Markov if and only if

$$\mathbb{P}[S_{t+1} | S_t] = \mathbb{P}[S_{t+1} | S_1, \dots, S_t]$$

- "The future is independent of the past, given the present" : $H_{1:t} \longrightarrow S_t \longrightarrow H_{t+1:\infty}$
- i.e. the state is sufficient statistic of the future.
- The environment state is Markov.
- The history is Markov.

→ nice case.
Full observability: agent directly observes environment state. $O_t = S_t^a = S_t^e$.

Formally, this is Markov decision process (MDP).

Partial observability: agent indirectly observes environment. $S_t^a \neq S_t^e$

Formally, this is partially observable MDP (POMDP).

Agent must create its own S_t^a , e.g.:

- Whole history
- Beliefs. $S_t^a = (P[S_t^e = s^1], \dots, P[S_t^e = s^n])$.
- RNN. $S_t^a = \phi(S_{t-1}^a W_s + O_t W_o)$.

Inside an RL agent

- Policy: agent's behaviour. $A = \pi(s)$ ↙ policy
- Value function: prediction of future reward
- Model: predicts what env. will do next.

Maybe less, maybe more.

Policy may be stochastic: $\pi(a|s) = P[A=a|S=s]$

$$V_\pi(s) = E_\pi [R_t + \gamma R_{t+1} + \gamma^2 R_{t+2} + \dots | S_t = s]$$

γ -discount reward (e.g. 0.99).

About model:

Transitions: P predict the next state (dynamics)

Rewards: R predicts the next (immediate) reward, e.g.

$$P_{ss'}^a = P[S'=s' | S=s, A=a]$$

$$R_s^a = E[R | S=s, A=a]$$

Categorizing RL agents: \swarrow implicit

- Value based: No policy, value function.
- Policy based: policy, no value function.
- Actor critic: policy & value function.
- Model free: policy and/or value function, no model.
- Model based: policy and/or val. funct, model.

Exploitation - exploration.

25.01.2023

L2. Markov Decision Processes

Almost all RL problems can be formalized as MDP

Def. A state S_t is Markov \Leftrightarrow

$$P[S_{t+1} | S_t] = P[S_{t+1} | S_1, \dots, S_t]$$

State transition prob.: $P_{SS'} = P[S_{t+1} = S' | S_t = S]$

State transition matrix:

$$P = \begin{bmatrix} P_{11} & \dots & P_{1n} \\ \vdots & & \\ P_{n1} & \dots & P_{nn} \end{bmatrix}$$

each row sums to 1.

Def. A Markov Process (or Markov Chain) is a tuple $\langle S, P \rangle$:

- S is a (finite) set of sets
- P is a prob. matrix

Markov Reward Process

Def. MRP is a tuple $\langle S, P, R, \gamma \rangle$:

- R - reward function, $R_s = \mathbb{E}[R_{t+1} | S_t = s]$
- γ - discount factor, $\gamma \in [0, 1]$

* give value to vertices \rightarrow goal

Def. The return G_t is the total discounted reward from time-step t :

$$G_t = R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1}$$

Def. The value function $v(s)$ of an MRP is the expected return starting from state s :

$$v(s) = \mathbb{E}[G_t | S_t = s]$$

Bellman Equation for MRPs:

- immediate reward R_{t+1}
- discounted value of successor state $\gamma v(S_{t+1})$

$$v(s) = \mathbb{E}[R_{t+1} + \gamma v(S_{t+1}) | S_t = s]$$

$$v(s) = R_s + \gamma \sum_{s' \in S} P_{ss'} v(s')$$

$$v = R + \gamma P v \quad (\text{in matrices})$$

$$\begin{pmatrix} v(1) \\ \vdots \\ v(n) \end{pmatrix} = \begin{pmatrix} R_1 \\ \vdots \\ R_n \end{pmatrix} + \gamma \begin{pmatrix} P_{11} & \dots & P_{1n} \\ \vdots & \ddots & \vdots \\ P_{n1} & \dots & P_{nn} \end{pmatrix} \begin{pmatrix} v(1) \\ \vdots \\ v(n) \end{pmatrix}$$

* n - number of states.

It can be solved directly:

$$v = (I - \gamma P)^{-1} R$$

$O(n^3)$ - comp. complexity.

Markov Decision Processes

Def. A MDP is a tuple $\langle S, \mathcal{A}, P, R, \gamma \rangle$:

- \mathcal{A} - a finite set of actions
- P is a state transition prob. matrix:

$$P_{ss'}^a = \mathbb{P}[S_{t+1} = s' \mid S_t = s, A_t = a]$$

- R is a reward function

$$R_s^a = \mathbb{E}[R_{t+1} \mid S_t = s, A_t = a]$$

Def. A policy π is a distribution over action, given states:

$$\pi(a|s) = P[A_t = a | S_t = s]$$

A policy defines the behaviour of an agent. There, policy depends only on s , not time i.e. policy are stationary: $\forall t > 0, A_t \sim \pi(\cdot | S_t)$

Def. A state-value function $V_\pi(s)$ of an MDP is the expected return starting from state s , and then following π

$$V_\pi(s) = \mathbb{E}_{\pi} [G_t | S_t = s]$$

Def. The action-value function $q_\pi(s, a)$

$$q_\pi(s, a) = \mathbb{E}_{\pi} [G_t | S_t = s, A_t = a]$$

expected return taking action a and following π from state s .


Bellman Expectation Equation:

$$V_\pi(s) = \mathbb{E} [R_{t+1} + \gamma V_\pi(S_{t+1}) | S_t = s]$$


$$q_\pi(s, a) = \mathbb{E}_{\pi} [R_{t+1} + \gamma q_\pi(S_{t+1}, A_{t+1}) | S_t = s, A_t = a]$$

$$V_{\pi}(s) = \sum_{a \in \mathcal{A}} \pi(a|s) q_{\pi}(s, a)$$

$$q_{\pi}(s, a) = R_s^a + \gamma \sum_{s' \in \mathcal{S}} P_{ss'}^a V_{\pi}(s')$$



$$V_{\pi}(s) = \sum_{a \in \mathcal{A}} \pi(a|s) \left(R_s^a + \gamma \sum_{s' \in \mathcal{S}} P_{ss'}^a V_{\pi}(s') \right)$$



$$q_{\pi}(s, a) = R_s^a + \gamma \sum_{s' \in \mathcal{S}} P_{ss'}^a \sum_{a' \in \mathcal{A}} \pi(a'|s') q_{\pi}(s', a')$$

In MDP :

$$V_{\pi} = R^{\pi} + \gamma \overset{\text{mean by } \pi}{P^{\pi}} V_{\pi}$$

$$V_{\pi} = (I - \gamma P^{\pi})^{-1} R^{\pi}$$

Def. The optimal state-value function $V_*(s)$ is the max. value function over all policies:

$$V_*(s) = \max_{\pi} V_{\pi}(s)$$

The optimal action-value function $q_*(s, a)$ is the max. action-value function over all π :

$$q_*(s, a) = \max_{\pi} q_{\pi}(s, a)$$

* it's max reward, not the best path.
↳ if it's known then we can find the best path.
↳ we should find q_* then go by increasing q_* path.

Def. $\pi \geq \pi'$ if $V_{\pi}(s) \geq V_{\pi'}(s), \forall s$

Theor. For any MDP

- There exists an optimal policy π_* that is better than or equal to all other policies, $\pi_* \geq \pi, \forall \pi$

- All optimal policies achieve the optimal value function, $V_{\pi_*}(s) = V_*(s)$


- All optimal policies achieve the optimal action-value function, $q_{\pi_*}(s, a) = q_*(s, a)$

$$\pi_*(a|s) = \begin{cases} 1, & \text{if } a = \operatorname{argmax}_{a \in \mathcal{A}} q_*(a, s) \\ 0, & \text{otherwise} \end{cases}$$


There's always optimal deterministic solution of MDP.

How do we find q_* ? By looking backwards.


Bellman Optimality Equation




$$V_*(s) = \max_a q_*(s, a)$$



$$q_*(s, a) = R_s^a + \gamma \sum_{s' \in S} P_{ss'}^a V_*(s')$$



$$V_*(s) = \max_a R_s^a + \gamma \sum_{s' \in S} P_{ss'}^a V_*(s')$$



$$q_*(s, a) = R_s^a + \sum_{s' \in S} P_{ss'}^a \max_{a'} q_*(s', a')$$

Bellman Optimality Equation is non-linear.
Many iterative solutions.

Extensions to MDPs

26.01.2023

L3. Planning by Dynamic Programming

Introduction

DP assumes full knowledge of the MDP.
It's used for planning in an MDP.

* Viterbi algorithm - what's

Policy Evaluation

Problem: evaluate policy π

Solution: iterative application of Bellman expectation backup

$$V_1 \longrightarrow V_2 \longrightarrow \dots \longrightarrow V_\pi$$

Using synchronous backups

$$V_{k+1}(s) = \sum_{a \in \mathcal{A}} \pi(a|s) \left(R_s^a + \gamma \sum_{s' \in \mathcal{S}} P_{ss'}^a V_k(s') \right)$$

$$V^{k+1} = R^\pi + \gamma P^\pi V^k$$

Policy Iteration

- Given a policy π

- Evaluate π : $V_\pi(s) = \mathbb{E}[R_{t+1} + \gamma R_{t+2} + \dots | S_t = s]$

- Improve π : $\pi' = \text{greedy}(V_\pi)$.

In general, need more iterations of improvement/evaluation
Always converges to π_* .

Modified policy iteration:

- Does policy evaluation need to converge to V_π ?
- Introduce stopping condition:
 - ϵ -convergence
 - k iterations.

Value iteration

Theor. (Principle of Optimality)

A policy $\pi(a|s)$ achieves the optimal value from state s , $V_\pi(s) = V_*(s) \iff$ for \forall state s' reachable from s , π achieves the optimal value from state s' , $V_\pi(s') = V_*(s')$

$$V_*(s) \leftarrow \max_{a \in \mathcal{A}} R_s^a + \gamma \sum_{s' \in \mathcal{S}} P_{ss'}^a V_*(s')$$

Intuition: start w/ final rewards and work backwards.

Problem: find optimal policy π

Solution: iterative application of Bellman optimality backup

$$V_1 \rightarrow V_2 \rightarrow \dots \rightarrow V_*$$

$$V_{k+1}(s) = \max_{a \in \mathcal{A}} \left(R_s^a + \gamma \sum_{s' \in \mathcal{S}} P_{ss'}^a V_k(s') \right)$$

in matrices: $V_{k+1} = \max_{a \in \mathcal{A}} \left(R^a + \gamma P^a V_k \right)$

Problem	Bellman Equation	Algorithm
Prediction	Bellman Expectation Equation	Iterative Policy Eval.
Control	Bellman Exp. Eq. + Greedy Policy Improvement	Policy Iteration
Control	Bellman Opt. Eq.	Value Iter.

Algos based on state-value function: $O(mn^2)$ per iter. $\rightarrow V_\pi(s), V_*(s)$

Algos based on action-value function: $O(m^2n^2)$ per iter. $\rightarrow Q_\pi(s,a), Q_*(s,a)$

Extensions to Dynamic Programming

Three simple ideas for asynchronous DP:

- 1) In-place dp
- 2) Prioritised sweeping
- 3) Real-time dp

1) In synchr. you save 2 versions of v .
In-place stores only one.

2) Use magnitude of Bellman error to guide state selection:

$$\left| \max_{a \in A} \left(R_s^a + \gamma \sum_{s' \in S} P_{ss'}^a v(s') \right) - v(s) \right|$$

Backup the state with largest remaining Bellman error

3) Only states that are relevant to agent
Use agent's XP to guide the selection of states
After each time step S_t, A_t, R_{t+1}

Backup S_t

$$v(S_t) \leftarrow \max_{a \in A} \left(R_{S_t}^a + \gamma \sum_{s' \in S} P_{S_t s'}^a v(s') \right)$$

Large DP suffers curse of dimensionality,
so we'll use sampling.

Contraction mapping

↳ Math is in lecture notes.

27.01.2023,

L4. Model-free prediction

Introduction

Estimate the value function of unknown MDP.
↳ policy evaluation

Monte-Carlo learning

Can only be applied to episodic MDPs:

- All episodes must terminate

Goal: learn V_{π} from episodes of XP under π

$$S_1, A_1, R_2, \dots, S_k \sim \pi$$

Return: $G_t = R_{t+1} + \gamma R_{t+2} + \dots + \gamma^{T-1} R_T$

Value function: $V_{\pi}(s) = \mathbb{E}_{\pi}[G_t | S_t = s]$ ←

Monte-Carlo uses empirical mean return.

$$V(s) \rightarrow V_{\pi}(s), \quad V(s) = \frac{S(s)}{N(s)}$$

as $N(s) \rightarrow \infty$

↑ sum of total return where s was visited
↑ number of times s is visited first time

First-visit-MC evaluation

Every-visit MC policy Estimation;

$N(s)$ increments every-time s encountered

$S(s) \leftarrow S(s) + G_t$ every-time s encountered

$$V(s) = S(s) / N(s)$$

The mean μ_1, μ_2, \dots of sequence x_1, x_2, \dots can be computed incrementally:

$$\mu_k = \frac{1}{k} \sum_{j=1}^k x_j = \mu_{k-1} + \frac{1}{k} (x_k - \mu_{k-1})$$

Incremental MC Updates

For each state S_t with return G_t

$$N(S_t) \leftarrow N(S_t) + 1$$

$$V(S_t) \leftarrow V(S_t) + \frac{1}{N(S_t)} (G_t - V(S_t))$$

In non-stationary problems, it can be useful to track running mean, i.e. forget old episodes:

$$V(S_t) \leftarrow V(S_t) + \alpha (G_t - V(S_t))$$

Temporal-Difference Learning

TD learns from incomplete episodes, by bootstrapping.

TD updates a guess towards a guess.

Simple TD learning algo TD(0)

- Update $V(S_t)$ toward estimated return
 $R_{t+1} + \gamma V(S_{t+1})$ (TD target)

$$V(S_t) \leftarrow V(S_t) + \alpha (R_{t+1} + \gamma V(S_{t+1}) - V(S_t))$$

TD target is biased estimate of $V_\pi(S_t)$
It has lower variance than the return

MC has high variance, zero bias

TD has low variance, some bias

MC converges to solution w/ minimum MSE

TD converges to the solution of max likelihood Markov Model.

→ exploits Markov property

→ doesn't exploit Markov property

n-Step return

$$G_t^{(n)} = R_{t+1} + \gamma R_{t+2} + \dots + \gamma^{n-1} R_{t+n} + \gamma^n V(S_{t+n})$$

$G_t^{(\infty)}$ - MC

$$V(S_t) \leftarrow V(S_t) + \alpha (G_t^{(n)} - V(S_t))$$

Averaging n-Step Returns

e.g. $\frac{1}{2} G^{(2)} + \frac{1}{2} G^{(4)}$

$$\frac{TD(\lambda)}$$

$$G_t^\lambda = (1-\lambda) \sum_{n=1}^{\infty} \lambda^{n-1} G_t^{(n)}$$

$$V(S_t) \leftarrow V(S_t) + \alpha (G_t^\lambda - V(S_t))$$

if terminal state, $\lambda^{\tau-t-1}$

Can only be computed for complete episodes.

$$\lambda = 1 - MC$$

$$\lambda = 0 - TD(0)$$

Eligibility traces

- Frequency heuristic: assign credit to most freq. state
- Recency heuristic: — / — / — / — / recent state

$$E_0(s) = 0$$

$$E_t(s) = \gamma \lambda E_{t-1}(s) + 1 \mathbb{I}(S_t = s)$$

Backward view TD(λ)

$$\delta_t = R_{t+1} + \gamma V(S_{t+1}) - V(S_t)$$

$$V(S) \leftarrow V(S) + \alpha \delta_t E_t(S)$$

When $\lambda=0 \Rightarrow E_t(S) = \mathbb{1}(S_t=S) = 1 \Rightarrow \text{TD}(0)$

$\lambda=1 \Rightarrow \text{MC}$

Theor. The sum of offline updates is identical for forward-view and backward-view TD(λ):

$$\sum_{t=1}^T \alpha \delta_t E_t(S) = \sum_{t=1}^T \alpha (G_t^\lambda - V(S_t)) \mathbb{1}(S_t=S)$$

27.02.2023

LS. Model-free control

Introduction

Model-free prediction: evaluates policy π
find best possible policy


Note: Try model easy football and try RL

On-policy: "learn on the job"

learn about π from XP using π

Off-policy: "look over someone's shoulder"

learn about policy π from XP sampled from $\mu \neq \pi$



Greedy Action Selection \rightarrow Bandit problem

Exploration problem


ϵ -greedy exploration:

$$\pi(a|s) = \begin{cases} \epsilon/m + 1 - \epsilon & \text{if } a^* = \underset{a \in \mathcal{A}}{\operatorname{argmax}} Q(s,a) \\ \epsilon/m & \text{otherwise} \end{cases}$$

→ Theorem \forall ϵ -greedy policy π , the ϵ -greedy π' with respect to Q_π is an improvement, $V_{\pi'}(s) \geq V_\pi(s)$.

policy improvement theorem

Idea: update V every episode:



Def. Greedy in the Limit with Infinite Exploration (GLIE):

- all state-action pairs are explored infinitely many times, $\lim_{k \rightarrow \infty} N_k(s, a) = \infty$

- the policy converges on a greedy policy, $\lim_{k \rightarrow \infty} \pi_k(a|s) = \mathbb{1}(a = \underset{a' \in \mathcal{A}}{\operatorname{argmax}} Q_k(s, a'))$


e.g. ϵ reduces to zero at $a_k = 1/k$.

GLIE MC Control

- Sample k^{th} episode using $\pi: \{S_1, A_1, R_2, \dots, S_T\} \sim \pi$
- $\forall S_t$ and A_t in the episode
$$N(S_t, A_t) \leftarrow N(S_t, A_t) + 1$$
$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \frac{1}{N(S_t, A_t)} (G_t - Q(S_t, A_t))$$
- Improve policy based on new action-value function
$$\epsilon \leftarrow 1/k$$
$$\pi \leftarrow \epsilon\text{-greedy}(Q)$$

Theorem. GLIE MC control converges to the optimal action-value function, $Q(S, a) \rightarrow Q_*(S, a)$

Updating Action-Value Functions w/ Saesa



$$Q(S, A) \leftarrow Q(S, A) + \alpha (R + \gamma Q(S', A') - Q(S, A))$$

Theor. Sarsa converges to the optimal action-value function, $Q(s,a) \rightarrow q_*(s,a)$ under the following conditions:

- GLIE sequences of policies $\pi_t(a|s)$
- Robbins-Monro sequences of step-sizes

$$\sum_{n=1}^{\infty} \alpha_t = \infty; \sum_{n=1}^{\infty} \alpha_t^2 < \infty. \quad \alpha_t:$$

n-Step Sarsa

$$q_t^{(n)} = R_{t+1} + \gamma R_{t+2} + \dots + \gamma^{n-1} R_{t+n} + \gamma^n Q(S_{t+n})$$

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha (q_t^{(n)} - Q(S_t, A_t))$$

Sarsa (λ)

$$q_t^{\lambda} = (1-\lambda) \sum_{n=1}^{\infty} \lambda^{n-1} q_t^{(n)}$$

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha (q_t^{\lambda} - Q(S_t, A_t))$$

forward-view (not online because wait till the end of an episode)

Backward View Sarsa (λ)

Use eligibility traces

$$E_0(s, a) = 0$$

$$E_t(s, a) = \gamma \lambda E_{t-1}(s, a) + \mathbb{1}(S_t = s, A_t = a)$$

$$\delta_t = R_{t+1} + \gamma Q(S_{t+1}, A_{t+1}) - Q(S_t, A_t)$$

$$Q(s, a) \leftarrow Q(s, a) + \alpha \delta_t E_t(s, a)$$

Off-policy learning

Evaluate $\pi(a|s)$ to compute $V_\pi(s)$ or $Q_\pi(s, a)$
while following $\mu(a|s)$.

$$\{S_1, A_1, R_1, \dots, S_T\} \sim \mu$$

Important because eg. see what people did
and do better by robot.

$$\begin{aligned} \mathbb{E}_{x \sim p}[f(x)] &= \sum P(x) f(x) = \\ &= \sum Q(x) \frac{P(x)}{Q(x)} f(x) \\ &= \mathbb{E}_{x \sim Q} \left[\frac{P(x)}{Q(x)} f(x) \right] \end{aligned}$$

Importance Sampling for Off-policy MC

$$G_t^{\pi/\mu} = \frac{\pi(A_t|S_t)}{\mu(A_t|S_t)} \frac{\pi(A_{t+1}|S_{t+1})}{\mu(A_{t+1}|S_{t+1})} \dots \frac{\pi(A_T|S_T)}{\mu(A_T|S_T)} G_t$$

$$V(S_t) \leftarrow V(S_t) + \alpha (G_t^{\pi/\mu} - V(S_t))$$

Very high variance, so don't work

Importance Sampling for Off-policy TD

$$V(S_t) \leftarrow V(S_t) + \alpha \left(\frac{\pi(A_t|S_t)}{\mu(A_t|S_t)} (R_{t+1} + \gamma V(S_{t+1})) - V(S_t) \right)$$

Much lower variance

Q-learning

$$A_{t+1} \sim \mu(\cdot | S_t)$$

$$A' \sim \pi(\cdot | S_t)$$

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha (R_{t+1} + \gamma Q(S_{t+1}, A') - Q(S_t, A_t))$$

$\exists \pi$ is greedy with respect to $Q(S, a)$

$\exists \mu$ is ϵ -greedy w.r.t. $Q(S, a)$

Then, $R_{t+1} + \gamma Q(S_{t+1}, A') = R_{t+1} + \max_{a'} \gamma Q(S_{t+1}, a')$

Q-learning Control Algorithm (SARSA MAX)

$$Q(S,A) \leftarrow Q(S,A) + \alpha (R + \gamma \max_{a'} Q(S',a') - Q(S,A))$$

Theor. Q-learning control ^{a'} converges to the optimal action-value function, $Q(S,a) \rightarrow Q_*(S,a)$

28.02.2023,

L6. Value Function Approximation.

Introduction

$$\hat{V}(s, w) \approx V_{\pi}(s)$$

$$\hat{Q}(s, a, w) \approx Q_{\pi}(s, a)$$

w-parameters (of NN, for example)

we'll consider approximators as • linear comb. of features

• NN

we require that methods are suitable for non-iid and non-stationary data.

Incremental Methods

Value Function Approx. using SGD

VFA Represent state by feature vector:

$$x(s) = \begin{pmatrix} x_1(s) \\ \vdots \\ x_n(s) \end{pmatrix}$$

$$\text{Linear VFA: } \hat{V}(s, w) = x(s)^T w = \sum_{j=1}^n x_j(s) w_j$$

$$\text{Update} = \alpha (V_{\pi}(s) - \hat{V}(s, w)) x(s)$$

$$= \text{step-size} \times \text{error} \times \text{feature value}$$

Table lookup is special case of linear VFA:

$$x^{\text{table}}(s) = \begin{pmatrix} \mathbb{1}(s=s_1) \\ \vdots \\ \mathbb{1}(s=s_n) \end{pmatrix}$$

Incremental Prediction Algorithms:

- For MC:

$$\Delta w = \alpha (G_t - \hat{v}(S_t, w)) \nabla_w \hat{v}(S_t, w)$$

- For TD(0):

$$\Delta w = \alpha (R_{t+1} + \gamma \hat{v}(S_{t+1}, w) - \hat{v}(S_t, w)) \nabla_w \hat{v}(S_t, w)$$

- For TD(λ):

Forward: $\Delta w = \alpha (G_t^\lambda - \hat{v}(S_t, w)) \nabla_w \hat{v}(S_t, w)$


Backward:

$$\delta_t = R_{t+1} + \gamma \hat{v}(S_{t+1}, w) - \hat{v}(S_t, w)$$

$$E_t = \gamma \lambda E_{t-1} + x(S_t)$$

$$\Delta w = \alpha \delta_t E_t$$

Control with VFA



Approximate policy evaluation

$$\hat{q}(\cdot, \cdot, w) \approx q_\pi$$

Action-value Function Approximation

AVFA

$$\hat{q}(S, A, w) \approx q_{\pi}(S, A)$$

$$\Delta w = \alpha (q_{\pi}(S, A) - \hat{q}(S, A, w)) \nabla_w \hat{q}(S, A, w)$$

feature: $x(S, A) = \begin{pmatrix} x_1(S, A) \\ \vdots \\ x_n(S, A) \end{pmatrix}$

linear AVFA: $\hat{q}(S, A, w) = x(S, A)^T w$

Incremental Control Algorithms:
 Same as prediction, but use $\hat{q}(S_t, A_t, w)$ instead of $\hat{v}(S_t, w)$

Mountain Car problem

Convergence of Prediction Algos

On/Off-Policy	Algo	Table	Linear	Non-linear
On	MC	✓	✓	✗
	TD(0)	✓	✓	✗
	TD(λ)	✓	✓	✗
	MC	✓	✓	✓
Off	TD(0)	✓	✗	✗
	TD(λ)	✓	✗	✗
Gradient TD	all	✓		

Convergence of Control Algos

Algo	Table	Linear	Non-Linear
MC	✓	(✓)	×
Sarsa	✓	(✓)	×
Q-learning	✓	×	×
Gradient Q-learning	✓	✓	×

(✓) = chatters around

Batch Methods

Gradient descent is not sample efficient

Batch try to find the best fitting value function
 ↙ experience

$$\mathcal{D} = \{ \langle S_1, V_1^\pi \rangle, \dots, \langle S_T, V_T^\pi \rangle \}$$

Least squares algo:

$$\begin{aligned} LS(w) &= \sum_{t=1}^T (V_t^\pi - \hat{V}(S_t, w))^2 = \\ &= \mathbb{E}_{\mathcal{D}} [(V^\pi - \hat{V}(S, w))^2] \longrightarrow \min \end{aligned}$$

SGD with Experience Replay

Repeat: 1. Sample state, value from XP

$$\langle S, V^\pi \rangle \sim \mathcal{D}$$

2. Apply SGD update:

$$\Delta w = \alpha (V^\pi - \hat{V}(S, w)) \nabla_w \hat{V}(S, w)$$

Converges to least squares solution:

$$w^\pi = \underset{w}{\operatorname{argmin}} \text{LS}(w)$$

Linear LS-prediction: \swarrow condition(?)

$$\alpha \sum_{t=1}^T X(S_t) (V_t^\pi - X(S_t)^T w) = 0;$$
$$w = \left(\sum_{t=1}^T X(S_t) X(S_t)^T \right)^{-1} \sum_{t=1}^T X(S_t) V_t^\pi$$

$$\text{LSMC: } V_t^\pi \approx G_t$$

$$\text{LSTD: } V_t^\pi \approx R_{t+1} + \gamma \hat{V}(S_{t+1}, w)$$

$$\text{LSTD}(\lambda): V_t^\pi \approx G_t^\lambda$$

Converges on- and off-policy both MC & TD.

LS Policy Iteration.

01.03.2023

L7. Policy Gradient Methods.

Introduction

$$\pi_{\theta}(s, a) = \mathbb{P}[a | s, \theta]$$

+ : policy methods can be stochastic
in partially observed env. (using features) we
lose Markov Property \Rightarrow no deterministic solution.

Goal: given $\pi_{\theta}(s, a)$ w/ θ , find best θ

To measure quality:

- start value: $J_1(\theta) = V^{\pi_{\theta}}(s_1) = \mathbb{E}_{\pi_{\theta}}(V_1)$

- average value: $J_{avV}(\theta) = \sum_s d^{\pi_{\theta}}(s) V^{\pi_{\theta}}(s)$

- average reward per time-step:

$$J_{avR}(\theta) = \sum_s d^{\pi_{\theta}}(s) \sum_a \pi_{\theta}(s, a) R_s^a$$

$d^{\pi_{\theta}}(s)$ — stationary distribution.

Find θ that maximizes J .

Finite Difference Policy Gradient

$$\Delta \theta = \alpha \nabla_{\theta} J(\theta)$$

$$\nabla_{\theta} J(\theta) = \begin{pmatrix} \frac{\partial J(\theta)}{\partial \theta_1} \\ \vdots \\ \frac{\partial J(\theta)}{\partial \theta_n} \end{pmatrix}; \quad \frac{\partial J(\theta)}{\partial \theta_k} \approx \frac{J(\theta + \epsilon \mathbf{e}_k) - J(\theta)}{\epsilon}$$

MC Policy Gradient

likelihood ratios: $\nabla_{\theta} \pi_{\theta}(s, a) = \pi_{\theta}(s, a) \nabla_{\theta} \log \pi_{\theta}(s, a)$
score function: $\nabla_{\theta} \log \pi_{\theta}(s, a)$.

softmax policy

$$\pi_{\theta}(s, a) \propto e^{\phi(s, a)^T \theta} \leftarrow \text{linear comb.}$$

$$\nabla_{\theta} \log \pi_{\theta}(s, a) = \phi(s, a) - \mathbb{E}_{\pi_{\theta}}[\phi(s, \cdot)]$$

gaussian policy

$$\mu(s) = \phi(s)^T \theta \leftarrow \text{can be parametrized}$$

$$a \sim \mathcal{N}(\mu(s), \sigma^2)$$

$$\nabla_{\theta} \log \pi_{\theta}(s, a) = \frac{(a - \mu(s)) \phi(s)}{\sigma^2}$$

$\phi(s)$
 \uparrow features
 $\phi(s, a)$

Theorem, For \forall diff. policy $\pi_\theta(s, a)$, for any of the policy obj. functions $J \in \{J_s, J_{avR}, \frac{1}{1-\gamma} J_{avV}\}$, the policy gradient is

$$\nabla_\theta J(\theta) = \mathbb{E}_{\pi_\theta} [\nabla_\theta \log \pi_\theta(s, a) Q^{\pi_\theta}(s, a)]$$

Algorithm idea: use return V_t as an unbiased sample of $Q^{\pi_\theta}(s_t, a_t)$

$$\Delta \theta_t = \alpha \nabla_\theta \log \pi_\theta(s_t, a_t) V_t$$

50:27: \leftarrow full algo.

Actor-Critic Policy Gradient

We use critic to estimate action-value function

$$Q_w(s, a) \approx Q^{\pi_\theta}(s, a)$$

critic params : w

actor params : θ

$$\nabla_\theta J(\theta) \approx \mathbb{E}_{\pi_\theta} [\nabla_\theta \log \pi_\theta(s, a) Q_w(s, a)]$$

$$\Delta \theta = \alpha \nabla_\theta \log \pi_\theta(s, a) Q_w(s, a)$$

The critic is solving policy evaluation problem.
We can do this using MC, TD, TD(0), LS.

[57:24] Algo

Theorem, (Compatible Function Approximation Theorem)

If the following conditions are satisfied:

① $\nabla_w Q_w(s, a) = \nabla_{\theta} \log \pi_{\theta}(s, a)$

② w minimizes MSE

$$\mathcal{E} = \mathbb{E}_{\pi_{\theta}} [(Q^{\pi_{\theta}}(s, a) - Q_w(s, a))^2]$$

Then the policy gradient is exact,

$$\nabla_{\theta} J(\theta) = \mathbb{E}_{\pi_{\theta}} [\nabla_{\theta} \log \pi_{\theta}(s, a) Q_w(s, a)]$$

Reducing Variance Using a Baseline

→ Reduce variance w/o changing expectation

We subtract baseline function $B(s)$

A good baseline $B(s) = V^{\pi_{\theta}}(s)$

advantage function → $A^{\pi_{\theta}}(s, a) = Q^{\pi_{\theta}}(s, a) - V^{\pi_{\theta}}(s)$

$$\nabla_{\theta} J(\theta) = \mathbb{E}_{\pi_{\theta}} [\nabla_{\theta} \log \pi_{\theta}(s, a) A^{\pi_{\theta}}(s, a)]$$

Estimating the Advantage Function

We can use TD error as unbiased estimate of the $A^{\pi_\theta}(s, a)$

$$\nabla_\theta J(\theta) = \mathbb{E}_{\pi_\theta} [\nabla_\theta \log \pi_\theta(s, a) \delta^{\pi_\theta}]$$

$\delta_v = r + \gamma V_v(s') - V_v(s)$
in practice, can use approx. TD error.

Policy Gradient w/ Eligibility Traces

$$\Delta \theta = \alpha (V_t^\gamma - V_v(s_t)) \nabla_\theta \log \pi_\theta(s_t, a_t)$$

$$\delta_v = r_{t+1} + \gamma V_v(s_{t+1}) - V_v(s_t)$$

$$e_{t+1} = \lambda e_t + \nabla_\theta \log \pi_\theta(s, a)$$

$$\Delta \theta = \alpha \delta_v e_t$$

Can be applied online, to incomplete sequences.

Natural Policy Gradient Fisher information matrix

$$\nabla_\theta^{\text{nat}} \pi_\theta(s, a) = G_\theta^{-1} \nabla_\theta \pi_\theta(s, a)$$

$$G_\theta = \mathbb{E}_{\pi_\theta} [\nabla_\theta \log \pi_\theta(s, a) \nabla_\theta \log \pi_\theta(s, a)^T]$$

05.03.2023.

L8. Integrating Learning and Planning.

Introduction

learn model from XP.

Model-Based Reinforcement Learning

- a model M is a representation of MDP $\langle S, A, P, R \rangle$, parametrized by η .
- assume that S and A are known
- $M = \langle P_\eta, R_\eta \rangle$ represent state transition & reward.
- assume independence:

$$P[S_{t+1}, R_{t+1} | S_t, A_t] = P[S_{t+1} | S_t, A_t] P[R_{t+1} | S_t, A_t]$$

Goal: Estimate M_η from XP $\{S_1, A_1, R_1, \dots, S_T\}$

$$S_1, A_1 \rightarrow R_1, S_2$$

...

$$S_{T-1}, A_{T-1} \rightarrow R_T, S_T$$

- $S, a \rightarrow r$ - regression problem
- $S, a \rightarrow S'$ - density estimation problem
- pick loss function
- find params η that minimizes loss

Table lookup model

$$P_{S,S'}^a = \frac{1}{N(S,a)} \sum_{t=1}^T \mathbb{1}(S_t, A_t, S_{t+1} = S, a, S')$$

$$R_S^a = \frac{1}{N(S,a)} \sum_{t=1}^T \mathbb{1}(S_t, A_t = S, a) R_t$$

Alternatively, record all $\langle S_t, A_t, R_{t+1}, S_{t+1} \rangle$ and then sample $\langle S, a, \cdot, \cdot \rangle$

Planning with a model:

- value iteration
- policy iteration
- tree search

Sample-Based planning

Use model to generate samples

$$S_{t+1} \sim P_{\gamma}(S_{t+1} | S_t, A_t)$$

$$R_{t+1} \sim R_{\gamma}(R_{t+1} | S_t, A_t)$$

Apply model-free RL: MC control, SARSA, Q-learning

→ Usually more efficient.

Integrated Architectures

We consider two sources of XP:

- Real : sampled from env.
- Simulated : sampled from model M

Dyna-architecture :

- learn and plan value function (and/or policy) from real and simulated XP.

! Dyna-Q algo ! 54:11 !

The changed environment is harder/easier
→ need exploration: Dyna-Q +
revisit states that haven't been visited for a while.

Simulation-Based Search

Forward Search : build search tree with the curr. s_t
use model of MDP to look ahead
just sub-MDP from now.

Simulate from now → apply model-free RL

Monte-Carlo Tree Search (Evaluation)

Given M, γ

Simulate: $\{ \underline{S}_t, A_t^k, R_{t+1}^k, S_{t+1}^k, \dots, S_T^k \}_{k=1}^K \sim M, \gamma, \pi$

Build search-tree

Evaluate states $Q(S, a)$:

$$Q(S, a) = \frac{1}{N(S, a)} \sum_{k=1}^K \sum_{u=t}^T \mathbb{1}(S_u, A_u = S, a) G_u \longrightarrow$$

$$\xrightarrow{P} q_{\pi}(S, a)$$

Select action: $a_t = \underset{a \in \mathcal{A}}{\operatorname{argmax}} Q(S_t, a)$

MC control applied to simulated XP

$$Q(S, A) \longrightarrow q_{*}(S, A)$$

maximin - the best way to act is to think that opponent is acting the best way too

TD search applies Sarsa to sub-MDP.

TD Search

For each step of simulation, update $Q(S,A)$ using Sarsa:

$$\Delta Q(S,A) = \alpha (R + \gamma Q(S',A') - Q(S,A))$$

select actions based on action values $Q(S,a)$.

e.g. ϵ -greedy

may also use function approx. for Q .

Dyna-2

stores:

- long-term memory
- short-term (working) memory

→ updated from real XP

← updated from sim. XP

10.03.2023

19. Exploration & Exploitation

Introduction

Approaches to exploration:

1. Random expl. (ϵ -greedy)
2. Optimism in the face of uncertainty
(prefer to explore uncertainty)
3. Information state space
(agent's information as part of its state)

State-action vs. Parameter exploration
pick diff. A each time S is visited
 $\pi(A|S, \underline{u})$ ←

We'll focus on state-action exploration.

Multi-Armed Bandits

tuple $\langle \mathcal{A}, \mathcal{R} \rangle$

\mathcal{A} -actions ("arms")

$\mathcal{R}^a(u) = P[R=u | A=a]$ - unknown

$A_t \in \mathcal{A}$ generates $R_t \sim \mathcal{R}^{A_t}$

$\sum_{t=1}^T R_t \longrightarrow \max$

Def. Action-value: $q(a) = \mathbb{E}[R|A=a]$

Def. Opt. value V_* : $V_* = q(a^*) = \max_{a \in \mathcal{A}} q(a)$

Def. Regret: $I_t = \mathbb{E}[V_* - q(A_t)]$

Def. Total regret: $L_t = \mathbb{E}\left[\sum_{j=1}^t V_* - q(A_j)\right]$

Count $N_t(a)$; Gap $\Delta_a = V_* - q(a)$

$$L_t = \sum_{a \in \mathcal{A}} \mathbb{E}[N_t(a)] \Delta_a$$

Good algos ensures small counts for large gaps.

Greedy can lock onto a subopt. solution.

Optimistic initialization: $Q(a) = r_{\max}$

Then act greedily: $A_t = \underset{a \in \mathcal{A}}{\operatorname{argmax}} Q_t(a)$

Optimistic greedy has linear total regret.

ϵ -greedy explores forever
has linear total regret.

pick a decay schedule for $\epsilon_1, \epsilon_2, \dots$

e.g.

$$c > 0$$

$$d = \min_{a | \Delta_a > 0} \Delta_a$$

$$\epsilon_t = \min\left\{1, \frac{c |d|}{d^2 t}\right\}$$

Decaying ϵ_t -greedy has logarithmic total regret.

Theorem (Lai & Robbins)

Asymptotic total regret is at least logarithmic in numbers of steps:

$$\lim_{t \rightarrow \infty} L_t \geq \log t \sum_{a | \Delta_a > 0} \frac{\Delta_a}{KL(R^a \| R^{a^*})}$$

Upper Confidence Bounds

Estimate upper conf. bound $U_t(a)$

$$A_t = \underset{a \in \mathcal{A}}{\operatorname{argmax}} Q_t(a) + U_t(a)$$

Theorem 1 (Hoeffding's Inequality)

Let X_1, \dots, X_t be i.i.d random variables in $[0, 1]$,
and let $\bar{X} = \frac{1}{t} \sum_{j=1}^t X_j$. Then:

$$\mathbb{P}[\mathbb{E}[X] > \bar{X}_t + u] \leq e^{-2tu^2}$$

Condition on selecting a :

$$\mathbb{P}[q(a) > Q_t(a) + U_t(a)] \leq e^{-2N_t(a)U_t(a)^2}$$

$$\text{Pick } p : e^{-2N_t(a)U_t(a)^2} = p$$

$$\text{Solve } U : U_t(a) = \sqrt{\frac{-\log p}{2N_t(a)}}$$

Reduce p , e.g. $p = t^{-4}$

Ensures we select opt. act. $t \rightarrow \infty$: $U_t(a) = \sqrt{\frac{2 \log t}{N_t(a)}}$

UCB1 algo: $A_t = \underset{a \in \mathcal{A}}{\operatorname{argmax}} Q_t(a) + \sqrt{\frac{2 \log t}{N_t(a)}}$

Theore. $\lim_{t \rightarrow \infty} L_t \leq 8 \log t \sum_{a | \Delta a > 0} \Delta a$

Bayesian Bandits

Exploits prior knowledge.

e.g. Gaussians

Probability matching

$\pi(a) = P[Q(a) = \max_{a'} Q(a') | R_1, \dots, R_{t-1}]$
 select action a according to $\pi(a)$

Thompson sampling:

$$\pi(a) = \mathbb{E}[\mathbb{1}(Q(a) = \max_{a'} Q(a')) | R_1, \dots, R_{t-1}]$$

Sample from posterior and select max.

Value of Information

Information State Space

\tilde{S} - summary of all information

Action A causes a transition to a new

\tilde{S}' with prob. $\tilde{P}_{\tilde{S}\tilde{S}'}^A$

MDP $\tilde{\mathcal{M}} = \langle \tilde{S}, \mathcal{A}, \tilde{P}, R, \gamma \rangle$

In Bernoulli case, MDP can be solved by DP.
The solution is Gittins Index.
→ e.g. drug problem (work / don't work).

Contextual Bandits

tuple $\langle \mathcal{A}, S, R \rangle$

MDPs

- for unknown or poorly est. state, replace reward function with r_{\max} .
- $\tilde{S} = \langle S, I \rangle$ accumulated information

L10. Classic Games.

best response $\pi_*^i(\pi^{-i})$ is optimal policy against all other opponents' fixed policies

Nash equilibrium: $\pi^i = \pi_*^i(\pi^{-i})$
→ every player's policy is a best response
→ if found, game is solved.

Nash equilibrium is fixed-point of self-play RL

For general games, Nash equil. isn't unique, but we'll look at classic games where it's unique.

Two-Player Zero-Sum Games

White & Black $R^1 + R^2 = 0$
↑ rewards

Methods of finding Nash eq.:

- Game tree search
- Self-play RL

Perfect information: all visible

Imperfect information: not all visible

Minimax Search

$$V_{\pi}(s) = \mathbb{E}_{\pi} [G_t | S_t = s]$$

$$V_*(s) = \max_{\pi_1} \min_{\pi_2} V_{\pi}(s)$$

minimax policy $\pi = \langle \pi_1, \pi_2 \rangle$ is a Nash eq.

instead, we use value function $v(s, w) \approx V_*(s)$

Chinook solved checkers in 2007

We can apply RL algorithms (MC, TD(0), TD(γ)) to games by making them self-play.

Policy improvement w/ afterstates

$$Q_*(s, a) = V_*(\text{succ}(s, a))$$

$$A_t = \underset{a}{\text{argmax}} V_*(\text{succ}(S_t, a)) \quad \text{for white}$$

$$A_t = \underset{a}{\text{argmin}} -1 - 1 - 1 - \quad \text{for black}$$

TD performed poorly in chess, checkers because of tactical nature of games.

TD Root

update value towards successor search value

$$V_+(S_t, w) = \underset{s \in \text{leaves}(S_t)}{\text{minimax}} v(s, w)$$

TD Leaf

update search value towards successor search value

TreeStrap

— / — / deeper — / —

Simulation-Based Search

UCT algo: MC + UCB algorithm

Smooth UCT Search
↳ Imperfect information games

